Fusion Plasma Ignition and Heating Systems

Kronos vs. Competitors

Table of Contents

Fusion Plasma Ignition and Heating Systems ***1***

Kronos vs. Competitors ***1***

Kronos's External Ignition & Heating Design **2**

Design Overview 2

Ignition Strategy 2

Maintenance Advantages 2

Heating Systems in Other Fusion Projects **4**

Conventional Tokamaks (ITER, CFS ARC, etc.) 6

Pulsed Magnetic Compression (Helion Energy) 7

Beam-Driven FRC (TAE Technologies) 8

Magnetized Target Fusion (General Fusion) 9

Inertial Confinement & Other Novel Approaches 11

Cost, Maintenance, and Commercial Outlook **13**

Capital Cost per MW of Heating 13

Maintenance and Operating Cost 15

Future Cost Trajectories 16

Kronos's External Ignition & Heating Design

Design Overview

The Kronos S.M.A.R.T. tokamak employs a unique external plasma heating and ignition system[1]. All major heating and current-drive components – including high-power neutral beam injectors (NBIs), radio-frequency (RF) generators, and even the "dual ohmic" induction coils used to initiate current – are located outside the vacuum vessel[2]. This means there is no bulky central solenoid or fragile internal heating structure inside the plasma chamber. Instead, Kronos drives its plasma current and temperature using externally-mounted systems that feed energy through ports in the reactor wall[3]. Crucially, these heating units are modular and can be shared across multiple reactor cores in Kronos's design[4]. In practice, Kronos envisions a power plant composed of several compact fusion cores all serviced by common external heating equipment. For example, a set of beam injectors or RF sources could service one core, then be redirected to another core as needed. This modular "hub-and-spoke" setup allows expensive heating hardware to be utilized more efficiently across reactors, potentially lowering the effective capital cost per megawatt of heating power delivered to each core[5].

Ignition Strategy

With an aneutronic fuel mix (initially deuterium–helium-3, or D–He^3), Kronos faces a higher ignition temperature threshold than conventional D–T reactors[6]. The reactor must reach ion temperatures on the order of ~50 keV (≈580 million °C) for robust D–He^3 fusion[7]. To achieve this, Kronos relies on its powerful external heating: high-energy neutral beams and RF waves that rapidly inject energy into the plasma, while the dual ohmic coils induce current and ohmic heating without a central solenoid[8]. This approach essentially "trades off easier ignition for a far superior overall power plant," accepting the need for strong external heating to ignite D–He^3[9]. Once the plasma is burning, some self-heating does occur – D–He^3 fusion produces 3.6 MeV alpha particles that remain confined and heat the plasma – but unlike a D–T reactor, Kronos does not rely on internal alpha particle heating alone to sustain the burn[10]. In fact, a significant fraction of the fusion energy in Kronos comes in charged protons (from the p–Li-6 reaction later in the fuel cycle) that are intentionally extracted for direct electricity conversion rather than being thermalized in the plasma[11]. This means external power injection continues to play a role in steady operation, topping up any energy shortfall to maintain the burn. Kronos's projected fusion gain factor is $Q \approx 40$, meaning the fusion power is about 40× the input power[12]. In practical terms, only on the order of ~2–3% of the reactor's output power needs to be fed back into the plasma via the external heaters[12]. In short, Kronos's ignition system is designed for high-output, steady-state operation, using external energy injection to ignite and sustain an aneutronic plasma that would not self-ignite as readily as a D–T plasma.

Maintenance Advantages

By keeping all heating hardware **outside** the neutron and plasma environment, Kronos makes routine maintenance dramatically easier[13]. There is no need to open the vacuum vessel for servicing injectors or replacing antennas – these can be accessed and swapped externally. This is in stark contrast to conventional tokamaks where key heating elements (e.g. internal coil windings or in-vessel RF antennae) reside inside the reactor and can only be repaired during lengthy shutdowns[14]. Kronos's first-wall panels and direct conversion structures **inside** the vessel are themselves modular and segmentable for quick replacement[15]. Likewise, all auxiliary systems like power supplies and fuel injectors are stationed outside for convenient access[16]. The result is a high uptime design: minimal downtime for maintenance, with rapid component swap-outs if a heating unit fails. In effect, Kronos's external ignition system functions more like industrial equipment than like delicate nuclear internals – technicians can service it in a straightforward manner, analogous to maintaining a turbine or boiler outside a fission reactor vessel. This should improve the

capacity factor and lower operation & maintenance costs over the plant's life^[1][2]. Kronos's low neutron output (only on the order of 1–2% of a D–T machine's neutron flux) further prolongs component life[17]. The external heating hardware avoids intense radiation damage entirely, and even the inner structures suffer far less activation and embrittlement than in D–T designs. All these features give Kronos a potential edge in maintenance frequency and cost: internal components might last years without replacement, and external systems are serviceable "on-the-fly."^[3]

Heating Systems in Other Fusion Projects

Fusion developers employ a variety of plasma heating and ignition methods, each with different cost and operational characteristics. Below, we compare Kronos's approach to the heating systems of other leading fusion projects, from conventional tokamaks to novel pulsed concepts.

Fusion approach	Heating / ignition driver	Initial cost (heating CAPEX)	Maintenance profile	Longevity / service interval
Kronos (aneutronic tokamak)	External neutral-beam injectors (NBI) + RF; plus external D–He^3 micro-fusion proton source feeding the core	NBI: ≈ **$2–5M per MW-beam** (ITER-class scaling); RF gyrotrons: ≈ **$2–3M per MW-RF** (turnkey w/ magnet) (AIP Publishing)	All heating hardware is **outside** the vessel → swap modules hot-side; low activation lowers servicing burden	External modules target **multi-year** life; periodic source/filament/neutralizer service similar to tokamak NBIs; RF tube life driven by cathode hours (AIP Publishing)
ITER/CFS ARC (D-T tokamaks)	Internal ohmic + **NBI (0.5–1 MeV)** + high-power RF (ECRH/ICRH)	NBI: ITER-like **tens of MW per injector**; cost basis lands around **multi-$10Ms per injector** (≈$2–5M/MW typical); RF gyrotrons ≈ **$2–3M/MW** (ResearchGate)	Beamline ion sources, neutralizers, cryopumps; in-vessel launchers see arcing/erosion; neutron fields complicate access → **remote handling** windows	Injector overhauls on **year-scale**; blanket/divertor change-outs on **multi-year** cycles (long outages)
Helion (pulsed FRC compression)	**Pulsed power & coils** compress and then directly recover energy (no steady NBIs/RF)	FOAK plant **$50–80M per MWe** (integrated, not a separate "heater" line-item); heavy pulsed-power stack CAPEX (energynews.pro)	No in-vessel heaters; main wear on **coils, switches, capacitors** at ~1 Hz; modules designed for field swaps	Component life measured in **millions of pulses**; replacements as preventive maintenance intervals accumulate (energynews.pro)
TAE (beam-driven FRC)	**Continuous NBIs** (10s of MW) sustain FRC; future p–B^{11}	Similar NBI/RF cost basis as tokamaks (tens of MW injector farms) → **$2–5M/MW (NBI)**,	External injectors; steady thermal loads (not pulsed); periodic ion-source,	Beam sources: **year-scale** component service; chamber parts at ends refreshed as needed

Fusion approach	Heating / ignition driver	Initial cost (heating CAPEX)	Maintenance profile	Longevity / service interval
		$2–3M/MW (RF) (AIP Publishing)	neutralizer, pump service	
General Fusion (MTF pistons)	**Mechanical pistons** implode a liquid-metal liner around magnetized target (pulsed)	Drivers are industrial (pistons, compressors) → **lower $/MW-driver** vs. NBIs; whole plant FOAK not public; prototype at **~1 shot/day** target rate initially (General Fusion)	Wear on **pistons/seals/valves** and liquid-metal systems; liner also serves as blanket/first wall	Service tied to **shot-rate**; frequent small component replacement acceptable if cheap & modular (General Fusion)
Zap Energy (SFS Z-pinch)	**Ohmic pinch** (MA-class current through plasma) stabilized by sheared flow (pulsed→aiming steady)	Capital light (capacitors + switches + coaxial chamber) → **very low $/MW-driver**; no large magnets/lasers (Taylor & Francis Online)	**Electrode erosion** and impurity control are dominant maintenance items; materials/devices under active R&D (AIP Publishing)	Longevity hinges on **electrode life** (goal: millions of shots) or moving to architectures minimizing solid contact (AIP Publishing)
First Light Fusion (projectile ICF)	**Hyper-velocity projectile** shocks target (pulsed); liquid-metal wall for energy capture/breeding	Company target: **≤$1B** for ~**150 MWe** pilot (integrated); per-shot target & launcher costs dominate OPEX (First Light Fusion)	Repeating **launcher barrel/valve** wear; target manufacture logistics; liquid-metal chamber absorbs debris	Service cadence driven by **gun life** (barrel/armature) and target feed system

Table 1 Heating System Economics (Table 1). Comparison of ignition/heating system capital costs, maintenance profile, longevity, and commercialization factors across leading fusion energy concepts, normalized to order-of-magnitude estimates for first-of-a-kind (FOAK) plants.

Conventional Tokamaks (ITER, CFS ARC, etc.)

Mainstream tokamaks using the deuterium–tritium (D–T) fuel cycle achieve ignition by a combination of internal and external heating. Initially, a large internal solenoid (central transformer) provides ohmic heating – driving a current through the plasma to heat it like a resistive element – and additional heating is supplied by NBIs and RF injectors mounted around the vessel[18]. For example, the ITER reactor will use two mega-scale neutral beam injectors delivering ~33 MW of 1 MeV D^0 beams, plus ~20 MW of high-frequency RF heating (ECRH/ICRH), to bring the plasma to ignition conditions[19][20]. These systems are proven and effective: past tokamaks like JET and TFTR in the 1980s–2000s all used tens of megawatts of neutral beams and RF to heat plasmas into the hundreds of millions of degrees[21]. Once D–T fusion reactions begin, the 3.5 MeV alpha particles produced in the plasma provide substantial self-heating. In a burning D–T plasma, this alpha heating can eventually dominate, making the plasma largely self-sustaining ("ignited") if confinement is good enough[22]. In essence, external heating is used to push the plasma up to ignition temperature, after which the internal alpha self-heating takes over most of the workload, reducing the recirculating power needs during steady burn.

However, this approach comes with heavy engineering overhead. The central solenoid and other in-vessel components in conventional tokamaks are bulky, complex, and subject to extreme neutron bombardment. They must be surrounded by a thick lithium blanket and neutron shield to absorb 14 MeV neutrons and breed tritium, which complicates maintenance. Replacing an internal component (e.g. a damaged coil or blanket module) entails a long shutdown with remote robotic handling; downtimes of many months for refurbishment are expected over a power plant's lifetime[23][24]. The capital cost of these heating systems is also high: neutral beam injectors are massive, expensive devices – ITER's injectors, for instance, are tens of meters long with cryogenic pumps and high-voltage accelerators[25]. While exact figures are proprietary, industry estimates put the cost of a large NBI on the order of **$2–5 million per MW** of beam power[26]. RF generator systems (gyrotrons, antennas, etc.) likewise cost on the order of ~$1 million per MW for high-frequency heating sources[27]. These systems add significantly to the up-front build cost of a fusion plant. In operation, the maintenance of such systems can be challenging – NBI beamlines have many delicate parts (ion sources, neutralizers) that require periodic replacement or cleaning, and RF antennas near the plasma can suffer arcing or erosion[28]. Traditional designs mitigate this by placing as much of the hardware as possible in accessible areas (for example, in ITER much of the injector hardware is in an external beam hall, with only a beam port penetrating the vessel)[29]. However, the central solenoid remains an Achilles' heel: it cannot be removed or repaired easily once the reactor is assembled[30]. Even advanced tokamak projects like Commonwealth Fusion Systems' ARC are working around this limitation by using demountable magnets to make internal components easier to swap, but they still must handle intense neutron loads and a complex internal blanket[31]. These factors contribute to projected first-of-a-kind costs in the billions for conventional D–T tokamak power plants[32]. (Indeed, CFS has a ~200 MW pilot plant contract with Google in the 2030s that is likely valued in the hundreds of millions, reflecting expected ~$10M+ per MW costs initially[33].) Over time, proponents believe mass-production and design simplification (e.g. modular construction and standardized blanket cartridges) will bring D–T tokamak costs down, but the need for neutron mitigation and large in-vessel structures sets a baseline of complexity and cost that aneutronic designs like Kronos seek to bypass[34].

Kronos's approach: By eliminating the need for a central solenoid through its dual external ohmic coils, Kronos avoids one of the major complexity and maintenance pain points of conventional tokamaks[30]. It forgoes the heavy internal structures and neutron-intensive blanket systems of a D–T tokamak, aiming to achieve ignition with external heating and then manage a steady burn without relying on an internal alpha-heating wave. Kronos's design thus seeks to bypass the entrenched cost drivers of conventional tokamaks (extensive shielding, tritium handling, frequent internal component replacement) in favor of a leaner approach focused on external, easily accessible systems.

Pulsed Magnetic Compression (Helion Energy)

Helion Energy's approach is radically different from a steady-state tokamak. It uses a pulsed, **non-ignition** fusion cycle to heat plasma, avoiding both the need for a sustained burning plasma and the use of massive steady-state heating systems[35][36]. In Helion's linear device, two field-reversed configuration (FRC) plasmoids are formed at opposite ends of a cylindrical chamber, accelerated towards each other, and then violently compressed together by a rapidly increasing magnetic field in the center[37][38]. The plasma heating occurs intrinsically during this compression – essentially converting kinetic and magnetic energy into thermal energy, much like a piston stroke compressing a gas[36]. This process can heat the merged plasma to over 100 million °C in a brief burst, triggering some D–He^3 fusion reactions in the fuel[39][40]. Notably, Helion does **not** attempt to ignite a self-sustaining burn; instead, as soon as the plasma begins to fuse and expand, Helion's coils start extracting electricity directly from the changing magnetic flux (per Faraday's Law) to recapture much of the input energy[41][42]. The plasma is then allowed to cool, and the cycle repeats at ~1 Hz (about one pulse per second). Helion describes this as a "non-ignition" pulsed system – it never relies on an internal alpha-heating burn wave, but rather on repetitive external energy injection each pulse, combined with extremely high-efficiency energy recovery[43]. In effect, Helion's magnet coils double as both the "heater" (during the compression phase) and the "generator" (during the expansion phase), eliminating the need for separate bulky heating hardware inside the reactor[44].

Cost and Efficiency: Helion's design aims to simplify the reactor and reduce cost by omitting cryogenic superconducting magnets and thermal conversion systems[45]. The pulsed copper magnets and power electronics *are* expensive, but they are more akin to industrial pulsed power equipment than to the custom, reactor-grade components in a tokamak[46]. Helion asserts this yields a smaller, cheaper device – they often note that their system "does not require the immense energy input of cryogenic superconducting magnets or large steam turbines"[47]. However, external analyses suggest the first-of-a-kind 50 MWe Helion plant (the **Orion** project for Microsoft) could still cost on the order of \$2.5–4 billion – i.e. roughly \$50–80 million per MW – when fully accounted[48][49]. This reflects the challenge that Helion must deliver and handle **huge** pulsed currents and forces (tens of MA of electric current, creating immense magnetic pressures) in a robust, repetitive system[50]. Over time, as they iterate designs and move to mass production, Helion hopes costs will fall drastically. Their ambitious long-term goal is on the order of only cents per kWh of electricity, and they plan to factory-build many small 50 MW units, leveraging learning-curve effects and bulk manufacturing[51]. This strategy is somewhat parallel to Kronos's modular multi-core concept (multiple smaller reactors built and fueled in parallel)[52].

Maintenance: In terms of maintenance, Helion's design benefits from having no fragile internal structure – the fusion happens in an empty cylindrical chamber between magnet coils, with no first-wall blanket or internal coil that must be periodically replaced[53]. There are no delicate components inside the plasma region at all. However, operating at ~1 Hz means Helion's pulsed magnets and electrical systems will experience significant cyclic stresses. Helion has demonstrated thousands of pulses on prototype devices, but scaling that to continuous 1 Hz operation for months or years will test the durability of components like capacitor banks, switching transistors, and coil insulation[54]. The company's approach to this is a modular one: because all the action happens in the magnets and electrical drivers, those can be **swapped out** relatively quickly if they degrade. One could imagine a Helion plant periodically doing an on-site swap of coil assemblies or power electronics modules, akin to servicing a large industrial electric machine[55]. Helion might also need to occasionally refurbish vacuum interface hardware or plasma injectors, but again, those are external to the fusion chamber and can be accessed without opening up a neutron-activated reactor vessel[56][57].

In summary, Helion's ignition/heating method is embedded in its pulsed fuel cycle – high-power magnetic compression replaces the continuous heating systems of a tokamak. It forgoes achieving a self-sustaining

plasma burn entirely, trading away the physics challenge of steady-state confinement in favor of a machine that "resets" after each pulse. This approach will be tested commercially soon: Helion's goal of delivering a 50 MW grid-connected fusion module by 2028 will demonstrate whether its unique pulsed heating cycle yields a cost-effective and reliable power plant[58]. Early estimates (like the ~$50M/MW figure above) suggest initial plants won't be cheap, but if Helion can rapidly iterate smaller units, it could drive costs down faster than more monolithic projects[59]. Kronos, notably, shares some of Helion's advantages – aneutronic fuel and direct electricity conversion – but it sticks to a steadier, ignited plasma model, whereas Helion essentially sidesteps steady-state plasma control by constantly pulsing the system[60].

Beam-Driven FRC (TAE Technologies)

TAE Technologies (formerly Tri Alpha Energy) is pursuing a steady-state **beam-driven** fusion reactor that also avoids D–T fuel. Their design uses a field-reversed configuration (FRC) plasma – essentially a self-confined torus of plasma without internal magnet coils – sustained by continuous injection of high-energy particle beams[61]. In TAE's current prototypes, large neutral beam injectors (on the order of 15–20 MW total) shoot energetic ions into the FRC to heat it and keep it stable[62][63]. This external injection is the primary heating and current-drive source, since an FRC has no central transformer to induce plasma current. TAE's ultimate goal is to fuse protons with boron-11 ($p–^{11}B$), an aneutronic reaction even harder to ignite than $D–He^3$; this means external heating may **always** be required to sustain the burn. In fact, analysis suggests an FRC plasma may never reach a classic ignition point where fusion alphas alone maintain the temperature – it might "always need some external heating" input to keep going[64]. TAE accepts this reality and is designing its system for high continuous power input and output: as long as the fusion power generated exceeds the beam power injected (i.e. $Q > 1$), the reactor can produce net energy[65]. In other words, TAE isn't striving for ignition in the traditional sense; it just needs a net power gain.

System Characteristics: TAE's reactors (such as the upcoming **Copernicus** and the planned **Da Vinci** pilot) will effectively be giant beam–plasma devices. They will likely use multiple beam injection ports around the machine, with powerful neutral beams (perhaps 20–50 MW total) providing both heating and particle injection to shape and sustain the plasma[66]. This is analogous to the auxiliary heating in tokamaks, except here it's not "auxiliary" – it is the primary means of maintaining the plasma. TAE's approach is steady-state (or very long pulse) rather than pulsed: the beams and plasma are on continuously, or at least for very long durations, which imposes steady thermal loads on components (less cyclic fatigue than a pulsed system, but more constant heat stress)[67]. The capital cost for TAE's heating and current-drive system will be substantial: they have developed proprietary high-energy beam injectors, but these involve large particle accelerators, high-capacity vacuum systems, and high-voltage power supplies similar in scale to tokamak NBI systems[68]. One advantage is that all this equipment sits outside the fusion core, behind shielding – maintaining TAE's injectors might be simpler than, say, replacing internal components of a D–T tokamak. The injectors can be reached without opening the reactor, much like Kronos's externalized system (in fact, Kronos's use of NBIs is inspired by approaches pioneered on tokamaks and by TAE)[69]. TAE's key differentiation is running the machine in steady-state: the beams and plasma are on nearly continuously (as opposed to pulsing), which reduces mechanical fatigue from thermal cycling but places a constant high heat load on components.

Cost and Outlook: TAE has raised over $1.2 billion in private funding and built a series of ever-larger FRC devices, demonstrating impressive control of FRC plasmas and recently achieving 100+ million °C temperatures in neutral-beam-heated plasmas (using D–D fuel)[70]. This progress indicates that high-power injection can indeed keep a high-beta FRC hot and stable. However, the efficiency of this scheme is a concern – the recirculating power fraction (the fraction of generated power used to run the injectors) might be significant if Q is not very high. TAE will need excellent beam coupling efficiency and *direct energy recovery* from fusion products to achieve a competitive net gain. They do plan to implement some form of

direct conversion or advanced energy capture for the charged alpha particles of p–^{11}B fusion, but that is challenging in their linear geometry[71]. Initially, they may operate with D–D or D–He^3 fuel mixtures (which produce some neutrons) and use a more conventional thermal cycle to convert that energy[72]. Thus, their first-generation plants could have substantial neutron shielding and look more like advanced beam-driven "neutron-lean" reactors rather than fully aneutronic ones.

In terms of maintenance, TAE's FRC has no internal magnet coils or solid blankets – much like Kronos and Helion, the design philosophy is to keep the vessel as simple as possible[73]. This should make it easier to swap out anything that does fail in the plasma chamber (for example, inner wall panels or beam entrance windows could be replaced relatively straightforwardly). TAE has published results suggesting they've improved electrode longevity and reduced impurities in their current device **Norman**, which bodes well for durability[74]. Still, a commercial TAE plant might have to replace certain components periodically – for example, beam sources degrade over time, and internal structures that are in the path of high-energy particles might need servicing. The good news is that these are generally outside the core confinement region (e.g., end electrodes or beam injectors), so they are maintainable without opening the main vessel frequently[75][76].

Comparing Kronos and TAE: In competitive terms, TAE's heating system is actually one of the most similar to Kronos's – both rely on a combination of neutral beams and other external systems to heat an aneutronic plasma. The key differences are in their physics approach and operational strategy. Kronos's tokamak approach seeks a high Q and a partly self-sustaining plasma (reducing the required continuous injection power), whereas TAE's FRC will likely operate at a lower Q (needing more beam power) but with a simpler, linear device geometry and possibly easier construction[77][78]. TAE's focus on p–^{11}B fuel is even more forward-looking – if it works, it produces no long-lived radioactive waste at all – but it means designing a machine that may never **ignite** in the traditional sense. Kronos, by contrast, plans to use D–He^3 fuel in its early years (a somewhat easier fusion fuel) to demonstrate ignition and net gain first, and only later transition to p–^{11}B fuel; essentially, Kronos hopes to prove out the concept on a more easily ignitable fuel, then gradually move to the hardest fuel, whereas TAE is "going for broke" on the hardest fuel from the start[79]. This strategic difference could affect timelines and cost: Kronos might achieve a net-energy D–He^3 pilot sooner (generating revenue or at least proving viability), whereas TAE may require more R&D time before a truly economic p–^{11}B reactor is ready[80]. In the long run, if both succeed, TAE's reactor could be very elegant – no neutrons, no Carnot-cycle thermal conversion (just beams in and electricity straight out). But it will likely **always** spend a larger fraction of its output to drive its own injectors compared to Kronos's more ignited approach[81]. In other words, an FRC might never reach the same level of internal self-heating as a tokamak, so more of the power it produces has to be diverted back into running the machine. Kronos is betting that by achieving some internal burn with a high-Q tokamak, it can minimize that recirculating power, whereas TAE is betting that it can make up for lower Q with a simpler design and no neutron-induced costs.

Magnetized Target Fusion (General Fusion)

General Fusion (GF) represents a very different philosophy: mechanical **compression** of a magnetized plasma target to achieve fusion. Their system injects a compact torus of plasma (a spheromak) into a spherical cavity filled with liquid metal, then uses an array of high-speed pistons to rapidly collapse the liquid metal liner around the plasma, squeezing it to fusion conditions[82][83]. In this approach, the "heating" of the plasma comes from external mechanical work – the kinetic energy of the moving liner is converted into heat and pressure on the fuel (similar in spirit to how a diesel engine's piston compresses and heats air)[84]. Like Helion, GF's process is pulsed: each shot physically compresses the plasma to fusion conditions (likely using D–T fuel initially), then the process repeats after reconditioning the

system[85]. Essentially, GF is attempting *magnetized inertial fusion* (or magnetized target fusion), using a mechanical driver instead of electromagnetic coils or lasers.

Cost Considerations: GF touts that its hardware is relatively low-cost and industrial: "the hardware (pistons, liquid metal) is low-cost and doesn't involve superconducting magnets or giant lasers"^[4]. Indeed, instead of millions of dollars in superconducting magnets or beam injectors, GF uses dozens of pneumatic or steam-driven pistons (a quite mature industrial technology) and a liquid lead–lithium mixture as the compression medium[86]. In principle, this could make the reactor vessel and driver hardware cheaper per unit energy than the intricate electromagnetic systems of other fusion approaches[87]. The liquid metal liner *doubles* as the heat transfer and neutron absorption medium – it directly captures the fusion energy as heat and also shields the solid structures from neutron damage[88]. This means GF can avoid expensive solid blankets and frequent replacement of first-wall structures; the fluid can be continuously circulated, cooled, and purified of any activation byproducts. Overall, if each compression cycle can be made reliable, the main operating cost would be the pumping power for the pistons and the maintenance of moving parts – which in classical engineering terms might be less expensive than maintaining large superconducting magnet systems[89].

Engineering Challenges: The flip side is that General Fusion must solve some very difficult engineering problems. Compressing plasma to fusion conditions requires an extremely symmetric and rapid liner implosion – any significant asymmetry or timing issue can quench the plasma before fusion occurs. Achieving this symmetry with mechanical pistons is nontrivial, especially shot after shot. Their current prototype will initially fire very infrequently (on the order of one shot per day) as they tune this process[90]. To become a viable power plant, they would need perhaps **>1 pulse per second**, which is a **huge** leap in repetition rate[91]. The wear on mechanical parts at that pulse rate is a major concern: seals, valves, and piston heads may fail under continuous high-frequency operation. GF's design mitigates some of this by using the liquid metal to soften the impact and by having no solid plasma-facing component that gets bombarded by each shot (the liquid liner reforms anew each time)[92]. But the pistons and compression chamber will still experience strong, repetitive stresses and some neutron irradiation (the liquid absorbs most neutrons, but not 100%)[93].

Maintenance Model: Maintenance for GF will likely involve periodically replacing or refurbishing the piston array and reprocessing the liquid metal (to remove accumulated radioactive byproducts or debris)[94][95]. On the positive side, those pistons are *accessible from outside* – like Kronos's external heaters, GF's drivers can be serviced without cutting open a vacuum vessel (each piston's actuator is outside the spherical compression chamber)[96]. Swapping out a piston unit could be as straightforward as a heavy machinery maintenance task, possibly accomplished in hours or days. General Fusion thus banks on using many **simpler, cheaper components replaced more often** instead of a few extremely expensive components that last longer[97]. It's a different economic model: many small "consumable" parts vs. a few giant high-tech parts. Which model is superior will depend on whether GF can prove high repetition with low-cost upkeep[98]. If, for instance, every piston has to be changed out every few weeks, that might still be economical if the pistons are cheap; but if they fail unpredictably or too often, it would hurt availability and drive up costs.

In sum, GF's heating system is an external **kinetic energy** driver – a mechanical compression system as opposed to Kronos's electromagnetic drivers. GF may have an advantage in capital cost per MW of driver power (steel pistons and compressors could cost less per MW of compression work than superconducting magnets or NBIs do per MW of heating)[99]. Moreover, GF's liquid-metal approach inherently comes with a built-in heat extraction loop (the liquid that gets hot from compression and neutron absorption can directly transfer that heat to a power conversion system)[100]. This bypasses the need for a separate coolant loop and potentially allows easier scaling of plant size (to get more power, one might simply use more pistons or a larger compression vessel)[101]. On the downside, ensuring precision and durability at scale is

unproven – as with other pulsed approaches, a lot relies on solving the repetition-rate and maintenance puzzle[102][103]. General Fusion hopes to demonstrate a breakeven single shot by ~2026 and then a pilot plant in the 2030s[103][104]. If they succeed, their fusion core might indeed be less costly than a tokamak core, but the balance of plant (e.g. turbines to convert heat to electricity) will still be needed unless they adopt something like Kronos's direct conversion. Since GF's output is mainly heat in the liquid metal liner, they are essentially building a fusion-powered heat source for a steam cycle, meaning the eventual electricity generation efficiency will be limited to ~40% (a typical steam turbine value) unless they incorporate advanced conversion methods[105]. Kronos avoids that limitation by directly converting charged particle energy to electricity, giving it a higher theoretical efficiency per reaction and reducing the size/cost of cooling systems needed[106].

Inertial Confinement & Other Novel Approaches

While Kronos and the companies above use magnetic confinement in various forms, a few notable projects use **inertial confinement** or other novel drivers to heat fusion fuel, quite unlike Kronos's magnetically-confined steady-state approach. For completeness, we consider these alternative heating systems and their commercial outlook briefly.

Laser ICF (e.g. NIF & Marvel Fusion)

The National Ignition Facility (NIF) in the U.S. proved in 2022 that 192 ultra-powerful laser beams delivering ~2 MJ to a tiny fuel capsule can **ignite** fusion in the target – a historic scientific result[107]. However, this was achieved with an enormous, one-of-a-kind machine. NIF's laser bank cost on the order of $3.5 billion to build and can fire at most a few shots per day[108], which is **orders of magnitude** away from what would be needed for a practical power plant. The energy yield per pulse was significant (NIF produced ~3 MJ of fusion energy in the December 2022 shot), but the cost per MJ of laser energy is extremely high and the laser efficiency is very low (only ~1% of the wall-plug electrical power is converted to laser light)[109]. For a commercial laser-fusion plant, radically new ultra-efficient drivers are needed to bring the cost per shot down. Companies like **Marvel Fusion** propose using next-generation high-repetition-rate lasers on p–^{11}B fuel targets, avoiding magnets entirely[110][111]. This would simplify the reactor design (essentially just a target chamber and banks of lasers). But at present, lasers are the most expensive way to deliver energy to a plasma per unit energy – NIF's example equates to on the order of $1 billion per MJ delivered, which is obviously untenable[109][112]. Marvel and others aim for breakthroughs like diode-pumped solid-state lasers to cut this cost by a factor of ~1000×[112]. Even if the cost and efficiency issues can be solved, there remains the question of chamber maintenance: each tiny fusion "micro-explosion" can damage the containing chamber's walls and the laser optics. Replacing final-focus lenses or cleaning debris out of the chamber might be needed frequently[113]. Some plans involve using renewable sacrificial layers (e.g. a flowing liquid wall or magnetic intervention) to shield components from direct exposure[114]. In summary, laser-based inertial fusion *can* achieve ignition (as shown by NIF's scientific demonstration), but "it's far from demonstrating a practical energy system" with current technology[115]. It likely remains a longer-term, high-risk pathway unless rapid advances in laser efficiency and chamber design occur.

Sheared-Flow Z-Pinch (Zap Energy)

Zap Energy is pursuing a fusion approach that uses a pulsed **electrical discharge** through plasma as the heating mechanism. In a Z-pinch, a powerful current is driven along a plasma column, which self-pinches (compresses) by magnetic forces, rapidly heating and compressing the plasma. Traditional Z-pinches are very prone to instabilities, but Zap adds a *sheared axial flow* to stabilize the plasma column. They have, in experiments, achieved stable Z-pinches lasting on the order of a few microseconds and produced significant D–D fusion neutron yields[116][117]. The appeal of this approach is its simplicity and low cost: essentially

it requires a coaxial tube (plasma gun) and capacitor banks – "a straight tube, two electrodes… cheap and modular," as Zap often says[118][119]. The heating power per pulse is large, but each Zap device is small and modular, and many could be run in parallel. The downside, similar to other pulsed systems, is **electrode erosion** and repetition rate limits. Zap reported firing 1,000 shots at ~0.1 Hz (one every 10 seconds) in recent tests[120], but scaling that up to continuous operation at, say, several Hz will be challenging. Each shot vaporizes a tiny bit of the electrodes, introducing impurities and gradually wearing the electrodes down[121][122]. Thus, while the *capital* cost per MW of pinch heating is very low (it's basically the cost of pulsed power electronics, likely far less than a neutral beam or laser system of equivalent power), the *maintenance* costs could be high if electrodes must be replaced often[123]. If Zap can find electrode materials or designs that last for millions of shots (or an automated way to swap electrodes frequently without much downtime), their system could become a very cost-effective fusion core – perhaps needing electrode swaps only after millions of shots (which could be automated)[124]. But until **Q ≈ 1** is demonstrated in the plasma and materials prove robust, it remains a high-risk, high-payoff concept[124][125].

Projectile Compression (First Light Fusion)

First Light Fusion uses a physical **projectile** fired at a target to generate a shockwave that compresses the fusion fuel – essentially a hybrid of inertial and magnetized target fusion[126]. In their concept, a high-velocity projectile (launched electrically or mechanically) provides the kinetic energy to rapidly compress and heat a plasma, similar in spirit to how lasers deliver energy in ICF, but potentially at much lower cost[127]. First Light has demonstrated very high shock pressures (around 3.6 TPa, or trillion pascals) in tests using the Sandia Z pulsed-power machine as a driver[128]. They are now planning a dedicated demo called **FLARE** that will combine an electromagnetic launcher with a projectile impact to implode fuel targets[129]. The promise here is cheap drivers (no expensive lasers or giant magnets – just EM railguns or even gas guns). A projectile can be mass-manufactured relatively inexpensively, and launch systems could perhaps fire many shots per minute if engineered properly[130][131]. Much like GF's pistons, however, the question is whether they can do it repetitively *and* capture the energy efficiently. Each shot will require a new target and projectile, and will create debris – i.e. ongoing consumable costs and wear on the chamber[132][133]. They avoid the large upfront cost of big magnet coils or giant lasers, so the *upfront* $/MW might be lower, but the operational costs per shot (fabricating targets, replacing sacrificial structures) could add up[134]. At present, First Light is still proving the physics; commercialization will hinge on developing robust, rapid-fire target technology and demonstrating sufficient gain per shot to be economical[135].

Overall Considerations: These alternative approaches underscore a common theme: there is a trade-off between physics and engineering in fusion heating systems. Designs like Kronos (and tokamaks in general) demand expensive, high-precision equipment (superconducting magnets, high-power NBIs/RF generators) to make the plasma physics easier – i.e. to confine the plasma steadily and get some self-heating once ignited[136]. On the other hand, pulsed and inertial approaches demand far less costly static infrastructure, but they shift the burden to **dynamic** systems and repetitive operation, which is an engineering challenge and can incur high ongoing costs (component fatigue, consumables, etc.)[137][138]. Kronos's external heating scheme is an attempt to capture the best of both worlds: by externalizing and modularizing its heating and power conversion systems, it aims for the reliability and efficiency of a steady-state tokamak **without** the usual maintenance nightmares[139]. At the same time, by choosing aneutronic fuel and direct conversion, Kronos eliminates many bulky systems required in D–T plants (no tritium breeding blankets, smaller radiation shields, no massive steam turbines)[140]. This should significantly reduce both capital and operating costs *if* the core physics can be made to work[141][142].

Cost, Maintenance, and Commercial Outlook

Capital Cost per MW of Heating

Company / Concept	Heating system type	Heating power needed (gross)	Effective heating CAPEX for 100 MWe	Notes
Kronos (aneutronic tokamak)	External NBIs (800 keV–1 MeV) + RF + laser-driven D–He^3 proton source	~60 MW external drive for Q ≈ 40	≈ **$120–180M** (NBI @ $2–3M/MW + RF, proton source ~$20M)	Shared injector hub can halve this if feeding multiple cores. Aneutronic → no tritium blanket costs.
ITER/CFS ARC (D–T tokamak)	NBIs (1 MeV) + RF + central solenoid	~50 MW heating to start burn, then alpha self-heating dominates	≈ **$200–300M** (ITER scaling: $4–6M/MW NBI + $2–3M/MW RF + solenoid)	Cost dominated by neutron shielding + tritium breeding integration. High maintenance outages.
Helion (pulsed FRC)	Pulsed magnetic compression (coils + pulsed power)	Entire plasma heated each 1 Hz cycle	FOAK integrated ≈ **$4–8B** for 100 MWe → effective $40–80M/MWe (driver included)	No "separate heater" line item — heating is the machine. ROI lever is energy recovery (~95%).
TAE (beam-driven FRC)	Continuous NBI farm (20–50 MW)	For 100 MWe net at modest Q (≈2–5), need 40–60 MW beams	≈ **$200–300M** (NBI cost basis $4–5M/MW)	Similar to tokamaks in injector costs, but simpler vessel. Always high recirculating beam load.
General Fusion (piston MTF)	Mechanical pistons compress liquid metal liner	Tens of MW per shot delivered as compression work	Hardware driver cost relatively low: ≈ **$50–100M** for 100 MWe module	Cost shifts to high OPEX (piston service). Conversion efficiency capped by steam cycle (~40%).
Zap Energy (sheared-flow Z-pinch)	Ohmic pinch (MA currents) from capacitor banks	Each pinch ~multi-MW equivalent; scale by rep rate	≈ **$20–40M** for driver stacks at 100 MWe	Lowest CAPEX per MW driver, but electrode erosion likely raises O&M significantly.
First Light Fusion (projectile ICF)	EM/gas gun projectile + target	≈ MJ-class shot energy at 1–10 Hz repetition	Company claims ~**$1B** pilot for 150 MWe → ~$670M/100 MWe	Heating tied to projectile launcher + mass-manufactured targets (OPEX driver).

Table 2 Heating Cost Benchmarks for 100 MWe Net Output (Table 2). Approximate capital expenditures required for heating or driver systems scaled to a 100 MWe net fusion plant, showing relative costs across different architectures. Estimates reflect FOAK-level deployments.

Among the various fusion approaches, there is a wide range in the **cost of delivering a megawatt of heating power** to the plasma. Traditional tokamaks tend to rank on the higher end – for instance, systems like ITER's 1 MeV neutral beams are first-of-a-kind technologies that cost many tens of millions for only tens of MW of output. Even smaller research tokamaks (JET, DIII-D, etc.) had neutral beam systems that likely cost several million dollars *per injector*. High-frequency microwave heating is also expensive, though somewhat more off-the-shelf (for example, a 1 MW gyrotron for ECRH costs on the order of ~$1 million including its power supplies)[143]. Therefore, a large D–T tokamak might easily spend on the order of \$100 million+ for ~50 MW of heating systems, not even counting the cost of the magnets and vessel[144]. Aneutronic tokamaks like Kronos will similarly require powerful heating sources – but Kronos plans to mitigate cost by sharing this equipment across multiple modules and by using as much commercially mature technology as possible[145]. For instance, Kronos can use "off-the-shelf" components adapted from existing fusion programs (e.g. neutral beam technology from tokamaks, high-power RF from particle accelerators)[146]. By designing the plant with many small tokamak cores served by a central set of heating units, Kronos could achieve a better economy of scale: instead of one reactor = one dedicated set of injectors, you might have one set of injectors servicing *N* reactors. If N = 4, for example, the cost per reactor for the heating system is roughly quartered (assuming the reactors can take turns or share power)[147][148]. This concept of a shared heating service is novel in fusion and is enabled by Kronos's modular layout. It is somewhat analogous to how a single central boiler can serve multiple processes in an industrial plant – one big investment instead of many smaller ones, yielding a lower cost per unit served[149]. Of course, this assumes the ability to distribute the heating effectively; Kronos will need very fast switching or multiple beamlines to feed several cores in parallel, and those implementation details will determine how much capital consolidation is possible in practice[150][151].

By contrast, pulsed approaches tend to have a lower initial cost per MW of input power, because they use common industrial hardware (capacitors, flywheels, pistons, etc.) rather than exotic high-field magnets or precision beamlines[152]. For example, Zap's entire pulsed power driver delivering a burst of a few MW is relatively cheap – basically just a bank of capacitors and switches. Helion's pulsed system, while not "cheap" in absolute terms, is basically big electrical equipment; their major costs are in power electronics and robust magnet coils. General Fusion's drivers are hydraulic pistons – again, something that scales in cost more linearly with size (doubling the compression power roughly doubles the piston system cost, with no extreme physics constraints)[153][154]. However, the **true cost per MW of useful heating** must factor in the efficiency and duty cycle of these systems. A neutral beam injector might convert ~40% of grid electrical power into beam power and run steady – you get what you pay for, continuously. A pulsed capacitor bank might dump energy with 90% efficiency, but if you can only fire once every 10 seconds, the *average* heating power is low unless you build many in parallel[155][156]. Helion's magnets compress the plasma with high efficiency and can fire every second, and crucially they recapture ~95% of the unused energy each pulse by design, so the net input per pulse is very low[157][158]. In that sense, Helion's **effective** cost per MW of heating could become very low *if* their energy recovery works – because they reuse the same input energy over and over, only adding a small amount each pulse to cover losses[159]. This is a unique advantage of their approach; they essentially lease the heating energy rather than spend it each time. Kronos and other steady-state systems don't have that cyclic energy recovery – energy put into heating is either transferred to the plasma or lost as heat, not returned. But since Kronos aims for a high Q (>>1), the fraction of output spent on heating is small. For example, at Q = 40, only about 2.5% of the reactor's output is used for heating, which might be acceptable even if that heating energy is "lost" to the plasma[160][161]. Helion's effective Q may be lower (because they deliberately don't self-ignite), but by recovering energy, they somewhat *cheat* the usual Q definition – they can report a high net gain even though each individual pulse isn't self-sustaining[162][163]. The bottom line is that different approaches distribute costs differently between upfront capital and ongoing operation, and a direct MW-to-MW cost comparison must account for efficiency and duty cycle.

Maintenance and Operating Cost

The frequency and complexity of maintenance strongly influence the economics of fusion power. Kronos's strategy to minimize maintenance downtime – via external, quick-swap components and low neutron damage – if successful, directly drives down operating costs. A fusion plant's revenue comes from uptime; every week offline for repairs is lost income. Traditional large tokamaks might require long maintenance outages every few years to replace neutron-battered blankets or divertor plates. For example, ITER expects to replace its blanket modules at least once in its 20-year lifetime, and likely to perform divertor replacements more often[164][165]. Each such operation can take months of remote handling work, translating to a high fixed O&M cost and a lower capacity factor. Kronos, with its "aneutronic" operation, avoids most of that: its first-wall modules are low-activation and can be replaced more like periodically scheduled maintenance (perhaps during an annual short outage) rather than massive emergency overhauls due to neutron damage[166][167]. Kronos's magnets, being superconducting, do need cryogenic support and careful monitoring, but by keeping neutron flux low and designing the magnet system with good thermal shielding, the magnets are expected to last many years without quenching or degradation[168]. This is a huge benefit – replacing a superconducting magnet (like those in ITER or ARC) would be a nightmare scenario due to cost and downtime. Kronos is engineered to avoid that need by protecting its magnets and using tough, radiation-resistant materials[169]. In short, Kronos aims for **high uptime** by minimizing in-vessel maintenance and keeping critical systems accessible.

Approach	Heating / driver CAPEX @ 100 MWe	Assumed CF (FOAK)	Effective CAPEX per average MWe (= $ / (100×CF))	Quick read
Zap (SFS Z-pinch)	**$20–40M**	70%	**$0.29–0.57M / avg MWe**	Cheapest driver if electrode life holds; O&M is the wild card.
General Fusion (piston MTF)	**$50–100M**	60%	**$0.83–1.67M / avg MWe**	Low driver CAPEX; repetition symmetry & piston wear drive O&M.
Kronos (aneutronic tokamak)	**$120–180M**	85%	**$1.41–2.12M / avg MWe**	External, modular NBIs + RF; sharing injectors across cores can lower this further.
TAE (beam-driven FRC)	**$200–300M**	80%	**$2.50–3.75M / avg MWe**	Beam farm CAPEX similar to tokamaks; steady beams → meaningful recirc.
ITER/CFS-style D-T tokamak	**$200–300M**	75%	**$2.67–4.00M / avg MWe**	High NBI/RF cost + neutron environment complicates maintenance access.
First Light (projectile ICF)	**$600–800M** (implied from ~$1B/150 MWe)	50%	**$12.0–16.0M / avg MWe**	Cost hinges on cheap targets + long-life launcher; steam-cycle efficiency cap.
Helion (pulsed FRC, integrated)	**$4–8B** (integrated FOAK)	60%	**$66.7–133.3M / avg MWe**	Heating = the machine; ROI bet is ~95% pulse energy recovery at scale.

Table 3 Capacity-Factor–Adjusted Heating Economics (Table 3). Effective cost per average MWe of heating/driver systems for a 100 MWe fusion plant, normalized by assumed FOAK

capacity factors. Highlights relative capital efficiency of ignition drivers across fusion approaches.

For the other approaches, maintenance requirements will differ. Helion's main maintenance costs will revolve around replacing electrical components (e.g. if a transistor module blows out or a coil fractures after many cycles). The nice part is that these are relatively modular – one could imagine a Helion plant swapping out coil assemblies or power electronics modules within a matter of days[170][171]. Helion might also need to occasionally refurbish vacuum interfaces or plasma injectors, but again these are external, so they can be accessed without breaking into the fusion chamber[56]. General Fusion will likely incur regular maintenance of its mechanical systems – pumps, valves, piston seals – similar to what is seen in conventional industrial plants with heavy machinery[57][172]. The cost of that is not negligible (mechanical systems need tending and parts replacement), but it's well-understood technology and typically scales with run hours. Importantly, GF's liquid-metal handling system will have to be maintained (impurities filtered out, leaks fixed, etc.), but liquid loops are a mature tech (similar to molten-salt loops in solar thermal plants, which are handled routinely in industry)[173][174]. TAE's beam-driven system will have a continuous power draw – so the electricity cost to run the injectors is a big operating expense until the reactor produces a large surplus. If their efficiency is low, the net electricity output might be small after powering the injectors, hurting economic viability[75]. They'll want to maximize Q or recapture energy from fast particles to minimize this recirculating power fraction. Maintenance for TAE will likely include replacing beam source components (filaments, ion sources), cryogenic pumps, etc., and possibly refurbishing components at the ends of the FRC that are struck by beams (e.g. end electrodes or beam target plates). Again, those parts are outside the core confinement region, so they are maintainable without opening the main vessel frequently[175][176]. Each approach has different O&M considerations, but designs like Kronos (and to some extent TAE and Helion) attempt to keep maintenance straightforward by putting as much as possible outside the high-radiation, hard-to-access zone.

Future Cost Trajectories

As fusion moves from first prototypes to commercial fleets, we expect significant cost improvements. Right now, estimates for *pilot plants* cluster in the tens of millions of dollars per MW of capacity (as noted above, on the order of ~$50M/MW for Helion's first unit, and similarly high for a first-of-kind tokamak pilot)[177][178]. This is clearly not competitive with existing power sources (for comparison, fission reactors are ~$6M/MW, and renewables even lower), but these are **first-of-a-kind** costs. Once designs are validated and production ramps up, modular designs like Kronos's or Helion's have a chance to undergo factory-style production and learning-curve cost reductions[179][180]. Kronos's concept of mass-producing small spherical tokamak cores and standardized external systems is aimed at driving the cost per unit down with volume. The shared-heating architecture further amplifies the benefit of scale – for example, one could imagine a 500 MW fusion plant made of 10 modules of 50 MW each, all fed by a set of injectors that are produced en masse. This begins to resemble an assembly-line product instead of a custom billion-dollar project built one at a time[181]. Helion, similarly, talks about eventually having factories turning out 50 MW fusion generators "like batteries." In one colorful claim, a Helion representative even mused about "a factory churning out twenty 50 MW reactors per day" in the far future[182]. While that was likely hyperbole, it illustrates the mindset: success in fusion will depend on scaling up manufacturing and scaling down unit cost. Aneutronic and low-neutron designs help in this regard because they can be more compact (no huge radiation shield volumes or remote-handling casks needed for activated components), meaning more modules can be built and transported easily[183].

On the other hand, if some approaches fail to tame their underlying physics challenges, their costs will remain high. For instance, if an aneutronic tokamak like Kronos had to significantly oversize its heating systems (say the achieved Q turns out to be only 5 instead of 40), it might need ~10× more beam power

than planned, making it more costly and complex – the very outcome Kronos is working to avoid by pushing for high field and high Q[184][185]. Similarly, if a pulsed concept like Zap Energy can't improve its electrode lifespan, the ongoing cost of replacing electrodes (or downtime to replace them) could negate the capital savings of the simple device[186][187]. It's notable that the fusion industry as a whole anticipates costs to drop as multiple prototypes come online in the 2030s. By that time, we'll have real data from several approaches: Commonwealth Fusion Systems hopes to demonstrate their ARC tokamak with net electricity by the early 2030s; Helion plans to have a 50 MW grid-connected plant by 2028; TAE aims for a prototype operating by ~2030 as well[188][189]. Kronos's roadmap is to have aneutronic pilot units in the 2030s too (leveraging partnerships and its fusion component factory)[190][191]. With multiple players coming to the table, competition should spur innovations that lower cost – whether it's better superconductors (reducing magnet cost per Tesla), improved efficiency in beam sources, cheaper power electronics, or novel materials that extend component life[192][193]. In the end, the fusion sector's commercial success will depend not just on *achieving* fusion in these different ways, but on making it economical and reliable through engineering refinement and industrial scaling.

Tables:

Table 1 Heating System Economics (Table 1). Comparison of ignition/heating system capital costs, maintenance profile, longevity, and commercialization factors across leading fusion energy concepts, normalized to order-of-magnitude estimates for first-of-a-kind (FOAK) plants.....5
Table 2 Heating Cost Benchmarks for 100 MWe Net Output (Table 2). Approximate capital expenditures required for heating or driver systems scaled to a 100 MWe net fusion plant, showing relative costs across different architectures. Estimates reflect FOAK-level deployments.....13
Table 3 Capacity-Factor–Adjusted Heating Economics (Table 3). Effective cost per average MWe of heating/driver systems for a 100 MWe fusion plant, normalized by assumed FOAK capacity factors. Highlights relative capital efficiency of ignition drivers across fusion approaches.....15

References:

1. Kronos Fusion Energy – *10-Year Strategy and Technology Analysis (Q3 2025).* Internal report (unpublished).
2. *Neutral-beam injection.* Wikipedia. (Describes principles and costs of NBI systems)[194][27].
3. S. Inoue et al., *"Operational Maintenance in ITER and Future Fusion Plants," Fusion Engineering and Design*, vol. 124, 2017. (Discussion of tokamak in-vessel component replacement)
4. General Fusion – *Technology Overview*. (General Fusion's description of its MTF approach and hardware costs)[86].
5. Helion Energy – *How to Make Fusion Electricity without Ignition.* (Helion blog post explaining its pulsed approach and energy recovery)[195][160].
6. M. Schroeder, *"Helion Energy defies physical and financial limits with its fusion power plant for Microsoft,"* **EnergyNews.pro**, Aug 2023. (Article on Helion's 50 MWe project and cost estimates)
7. World Nuclear News – *"Helion begins work on fusion power plant for 2028,"* Nov 2023. (Helion timeline and factory concept)
8. Kronos Fusion Energy – *"A Global Call for Fusion Energy Innovation and Partnerships,"* Jul 29, 2024 press release. (Kronos's partnerships and manufacturing strategy)
9. Kronos Fusion Energy – *"S.M.A.R.T. Fusion Energy Industrial Heating Unit,"* Dec 2024. (Kronos's external heating module for industrial applications)
10. TAE Technologies – Press Release, Jan 2023: "TAE achieves 100-million-degree plasma temp," (Milestone announcement on beam-heated plasma)
11. General Fusion – *"Preparing for Fusion Demonstration,"* General Fusion blog, Mar 2023. (Prototype plans and timeline for breakeven shot by 2026)
12. Marvel Fusion – *Technology Backgrounder*, 2023. (Laser-driven p–B^{11} approach and challenges in laser efficiency)
13. Zap Energy – *Technical Update*, 2023. (Results of 1,000-shot test and electrode erosion mitigation strategies)